



Dealing with system failure

HYSTRIX, NETFLIX OSS, SPRING CLOUD AND OTHER GEMS

-STUART INGRAM
@ARETHOSECLAMS

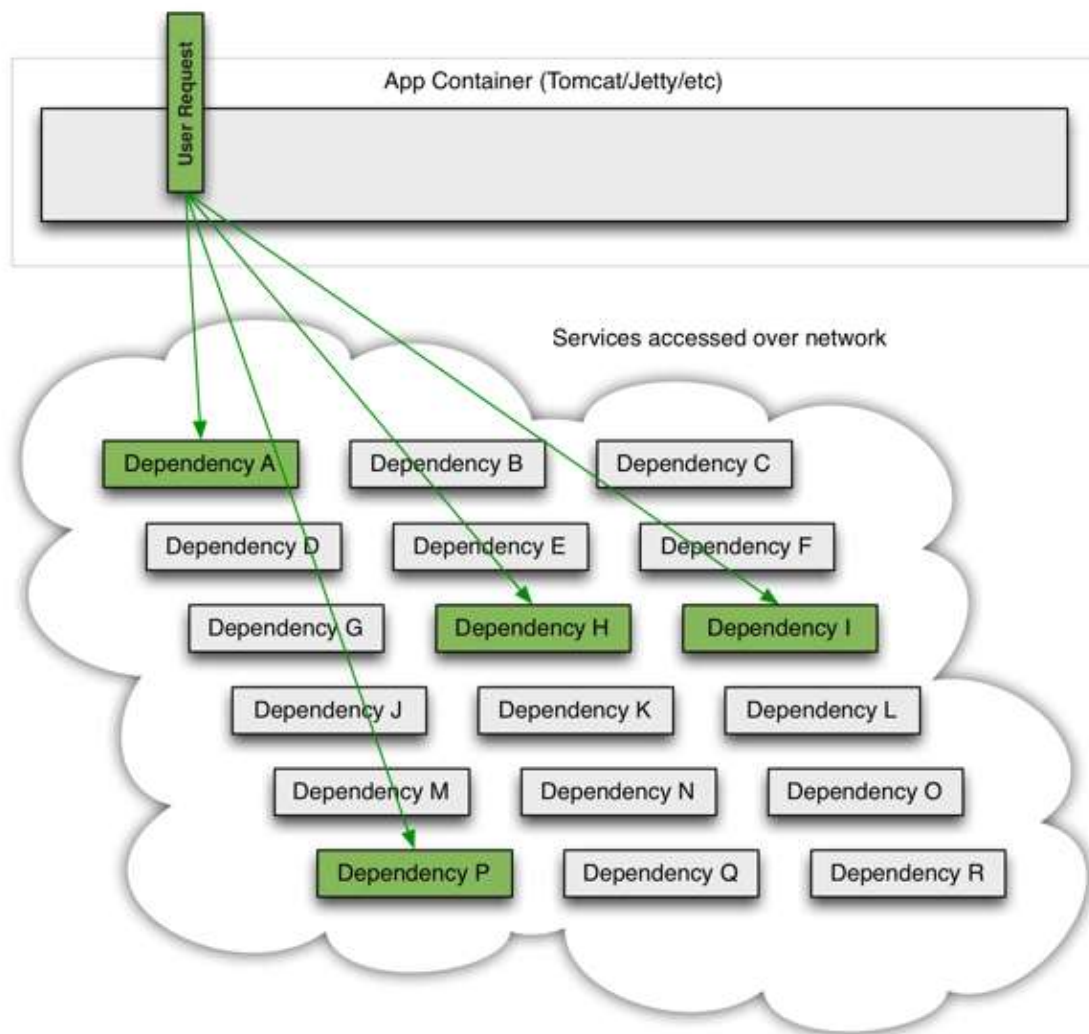
Expectations

- ▶ Problem Overview
- ▶ Circuit breaker pattern
- ▶ Concrete example - Hystrix
- ▶ Live demo
- ▶ Mystery bag

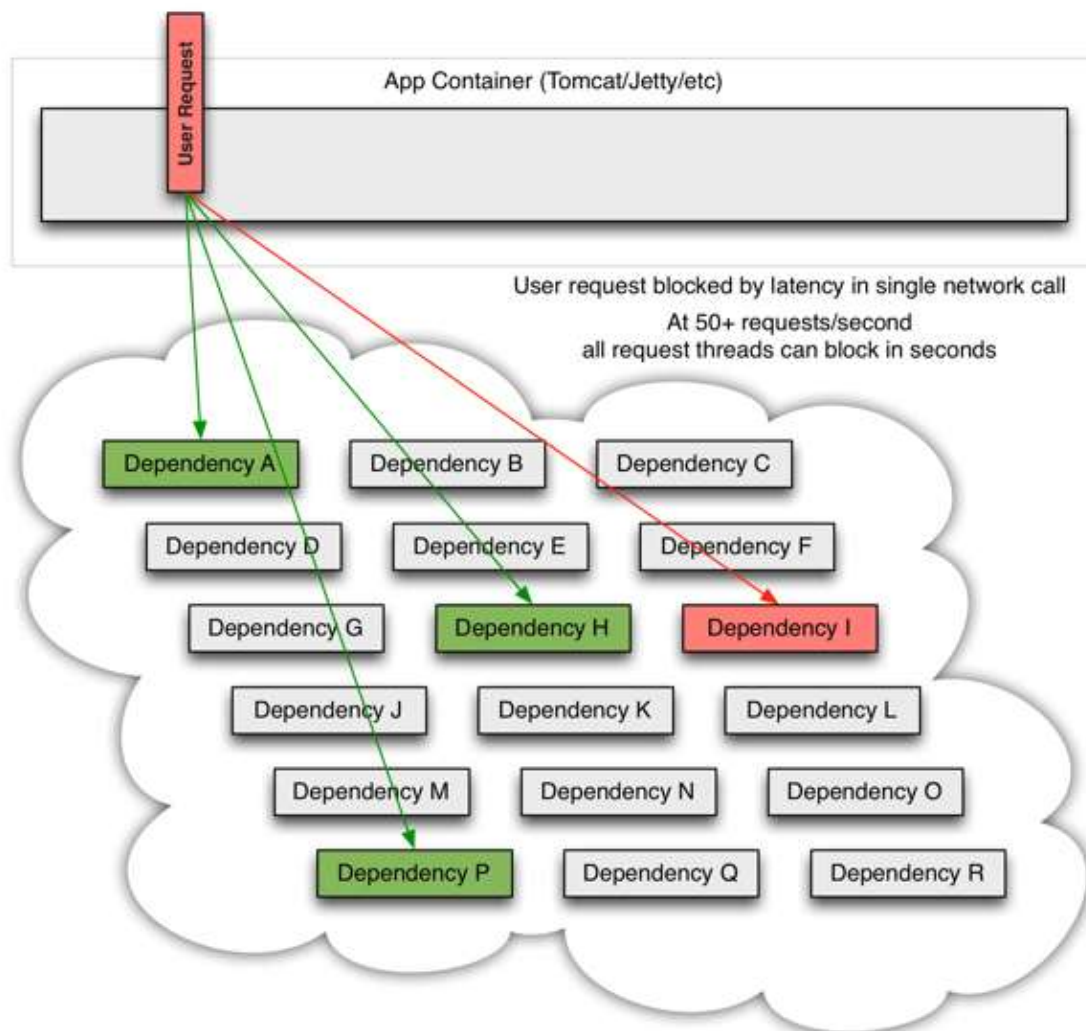
Failure

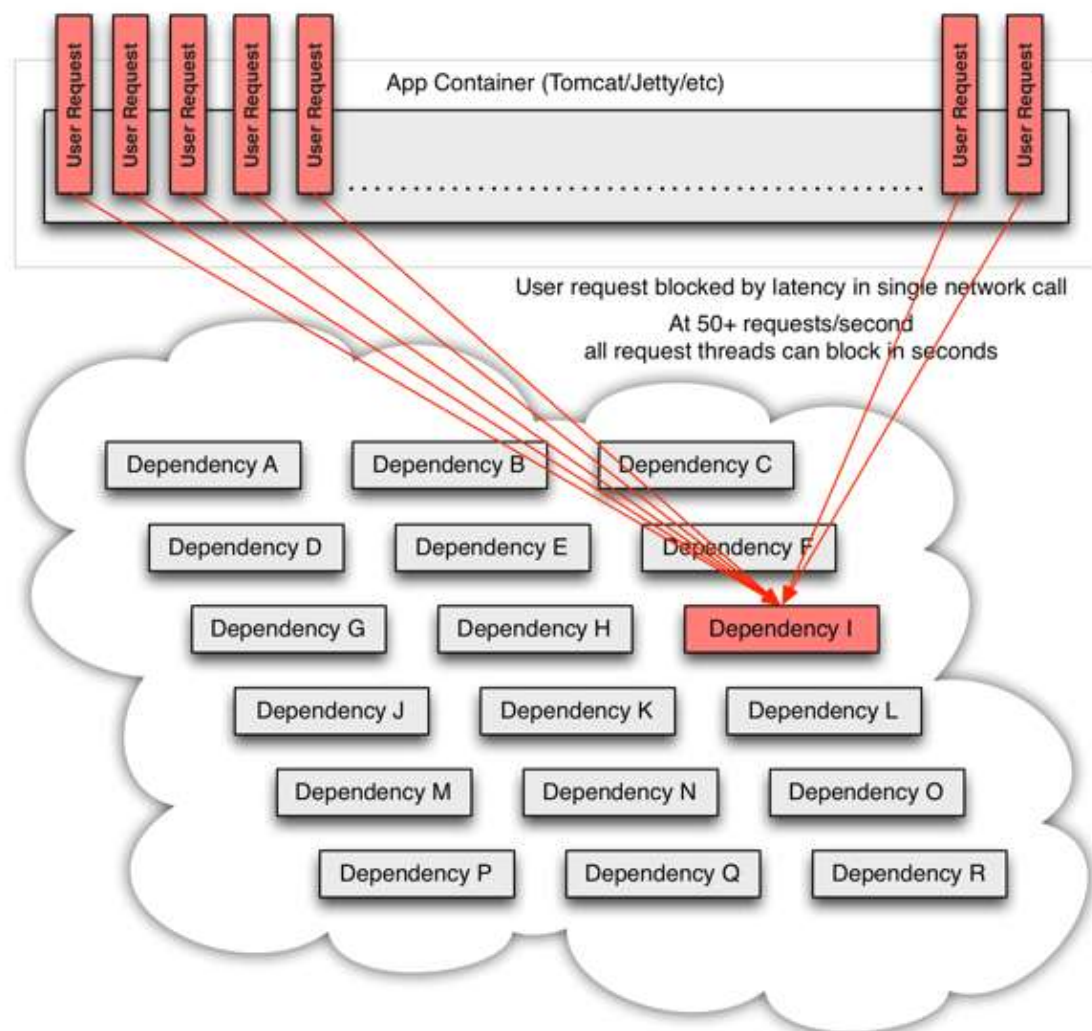
- 1 a : omission of occurrence or performance; specifically : a failing to perform a duty or expected action <failure to pay the rent on time>
b (1) : a state of inability to perform a normal function <kidney failure>
b (2) : an abrupt cessation of normal functioning <a power failure>
c : a fracturing or giving way under stress <structural failure>


- Merriam-Webster



<https://github.com/Netflix/Hystrix/wiki>







“*It's not what happens to you, but how you react to it that matters.*”

EPICTETUS – PHILOSOPHER ~ 55AD



Resilience

1 : the capability of a strained body to recover its size and shape after deformation caused especially by compressive stress

2 : an ability to recover from or adjust easily to misfortune or change

- Merriam-Webster

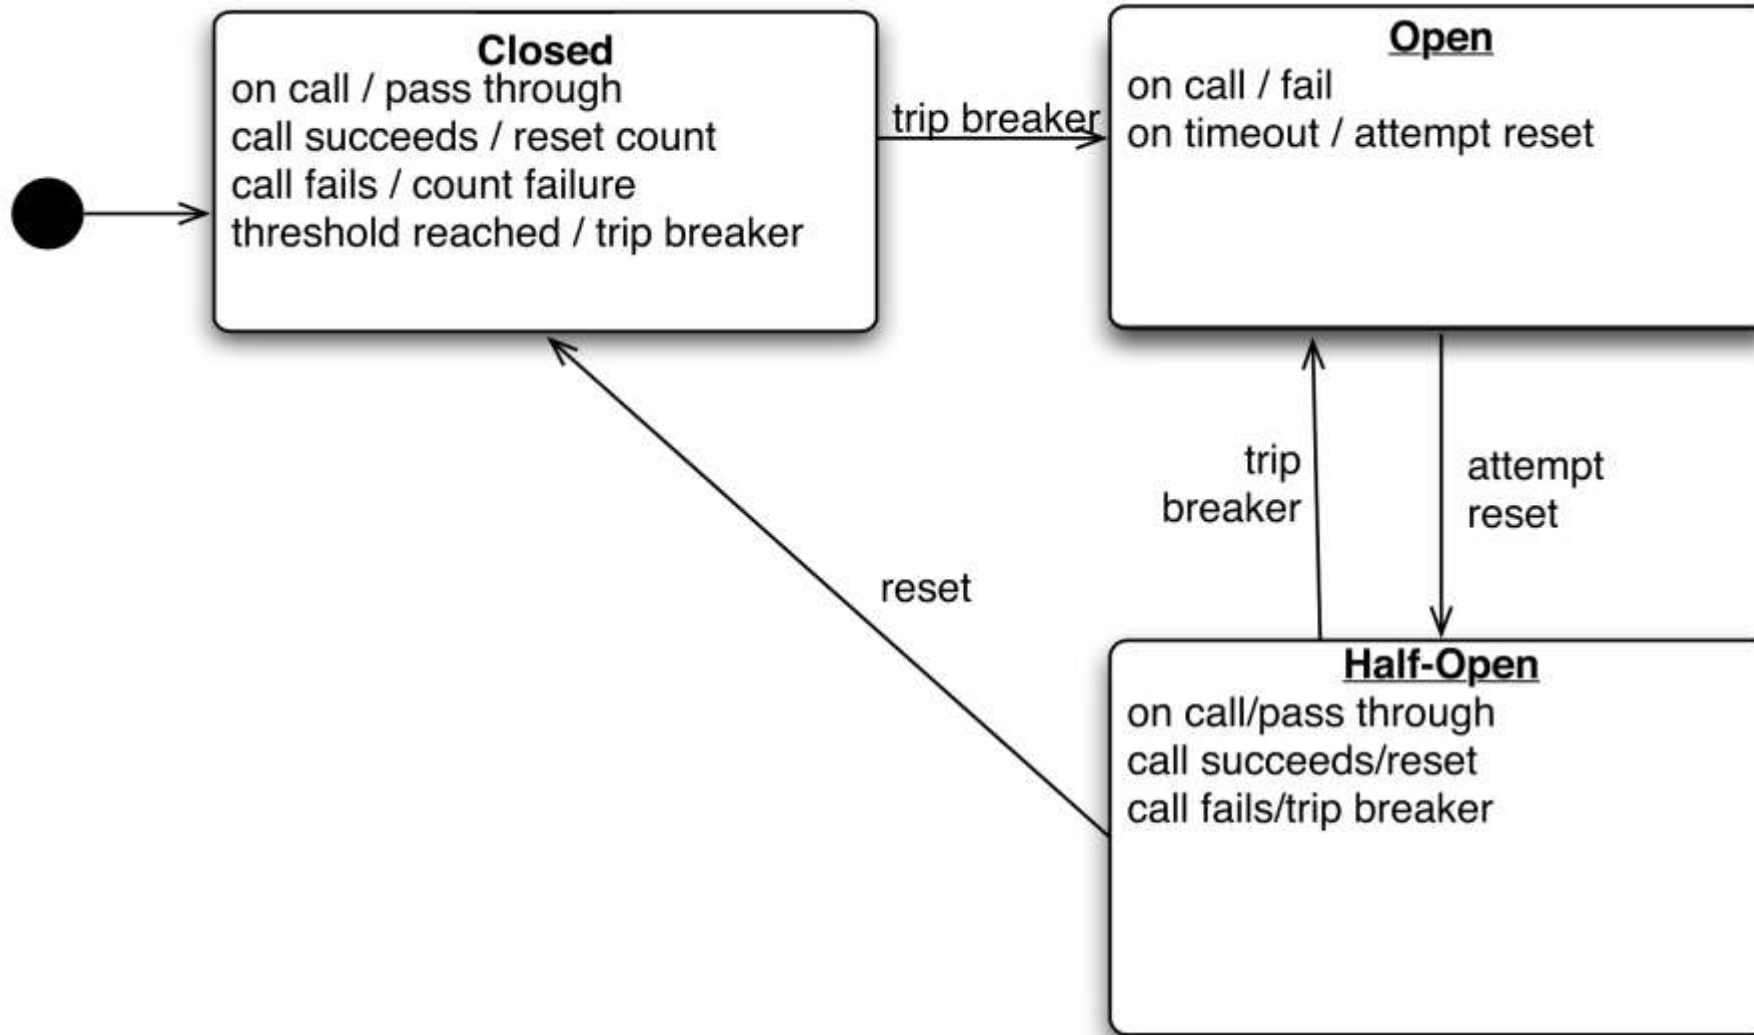
Circuit Breaker – Design Pattern

“A **circuit breaker** is an automatically operated electrical switch designed to protect an electrical **circuit** from damage caused by Overcurrent/overload or short **circuit**. Its basic function is to interrupt current flow after Protective relays detect faults condition.”

Wikipedia

"Release It! Design and Deploy Production-Ready Software" (2007)

Michael Nygard



Basic Properties

- ▶ Pros
 - ▶ Fail fast/ (cascading) fallback
- ▶ Cons
 - ▶ Overhead
- ▶ Trip properties
 - ▶ Latency
 - ▶ Absolute failure count
 - ▶ Percentage threshold
- ▶ Recovery
 - ▶ Retry window

Netflix OSS

- ▶ 1998 US DVD, 2007 US streaming, 2010 Canada & beyond.
- ▶ 75m subscribers. 190 countries. 125m hrs daily.
- ▶ Accounts for 36% of all downstream traffic in North America at primetime.
- ▶ Completed 8 year transition to AWS with strong micro service arch
 - ▶ Cloud architecture focuses heavily on failure & adaptability
 - ▶ Strong proponent of OSS with numerous contributions (203 on github)
 - ▶ Chaos Monkey (Simian Army – Chaos Kong)

Spring Cloud

- ▶ Spring (1.0.0 released 2002)
 - ▶ IoC
 - ▶ Source of all things XML
- ▶ Spring Boot (1.0.0 April 2014)
 - ▶ Convention over configuration
 - ▶ Annotation heavy
- ▶ Spring Cloud Netflix (1.0.0 March 2015)
 - ▶ Simple integration of the Netflix OSS Microservice stack
 - ▶ Swappable components

Hystrix – Netflix's Circuit Breaker

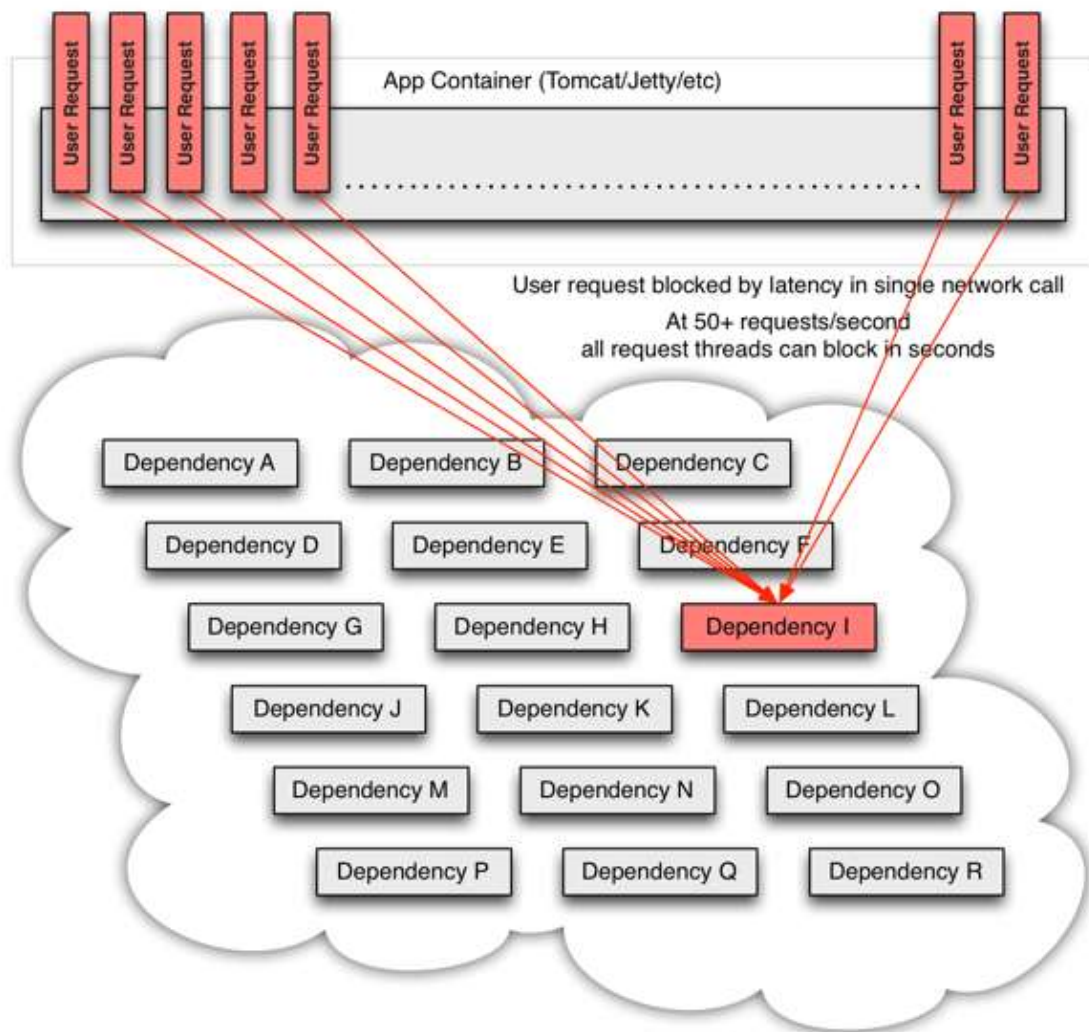
- ▶ Announced 2011
<http://techblog.netflix.com/2011/12/making-netflix-api-more-resilient.html>
- ▶ Library facilitating control of the interactions between these distributed services by **adding latency tolerance, fault tolerance logic** and *near realtime operational insights*.
- ▶ Tens of billions of thread-isolated calls are executed via Hystrix every day at Netflix.
- ▶ Hystrix library, Javanica, Spring Cloud Hystrix Starter

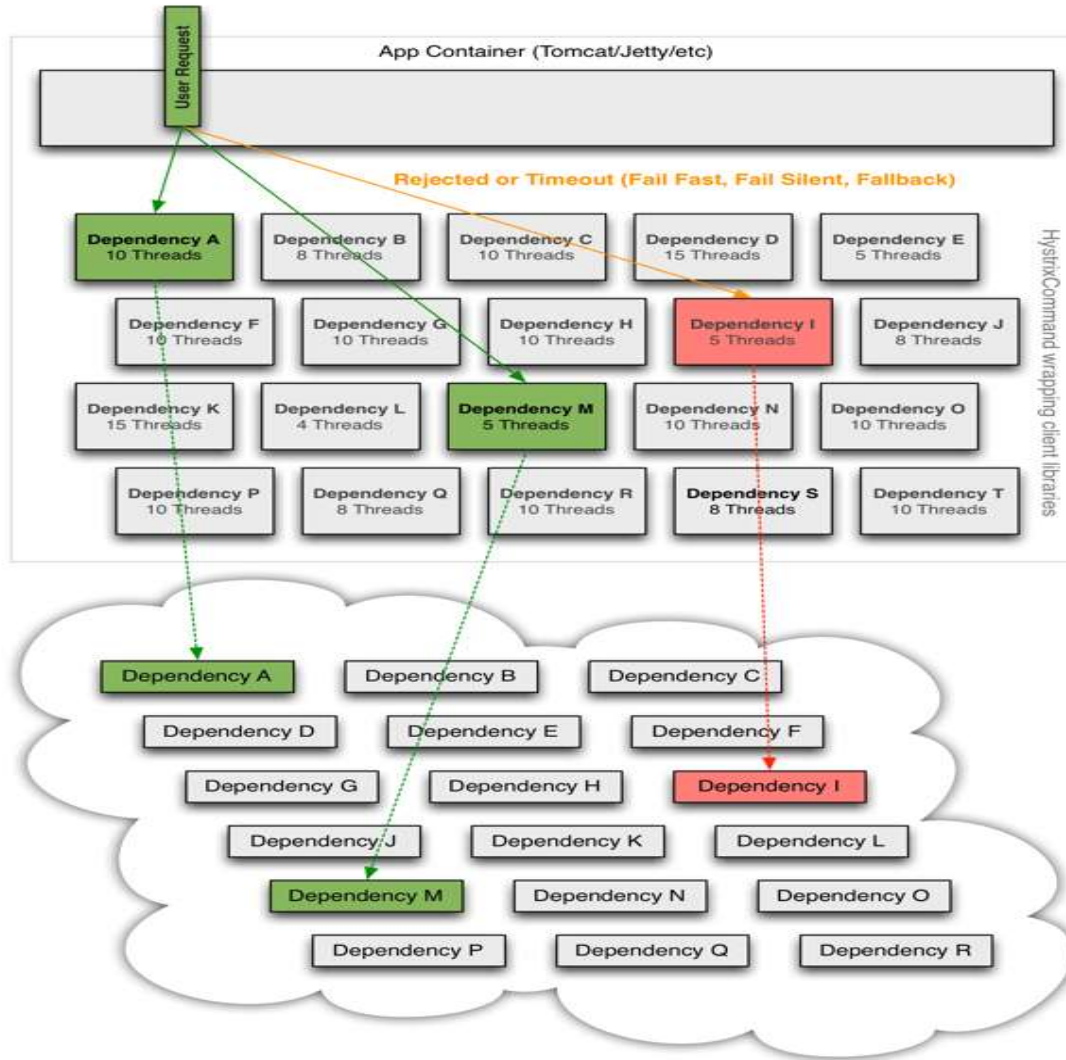
Configuration

- ▶ Circuit breaker (interaction)
 - ▶ Timeout
 - ▶ Fallback
 - ▶ Absolute error threshold
 - ▶ Percentage error threshold
 - ▶ Retry window
- ▶ Global/ per interaction
- ▶ Static / Dynamic

Configuration (cont.)

- ▶ Metrics (& Collapser)
 - ▶ Rolling window size
- ▶ Thread pool (external resource)
 - ▶ Size (core and queue)





Thread pools (Bulkheads)

- ▶ Shared by multiple commands/actions
 - ▶ Sensible defaults
 - ▶ GroupKey -> Class name (e.g. PaymentServiceGateway)
 - ▶ CommandKey -> Method name (e.g. depositCash)
 - ▶ Not configurable on a per-command basis without separate pool
- ▶ Resource limiting
- ▶ Isolation strategies

Usage

```
@HystrixCommand
```

```
@HystrixCommand(groupKey="UserGroup", commandKey =  
"GetUserByIdCommand")
```

```
@HystrixCommand(fallback='planBMethod')
```

- ▶ Synchronous
- ▶ Asynchronous (Futures)
- ▶ Reactive (Observable)

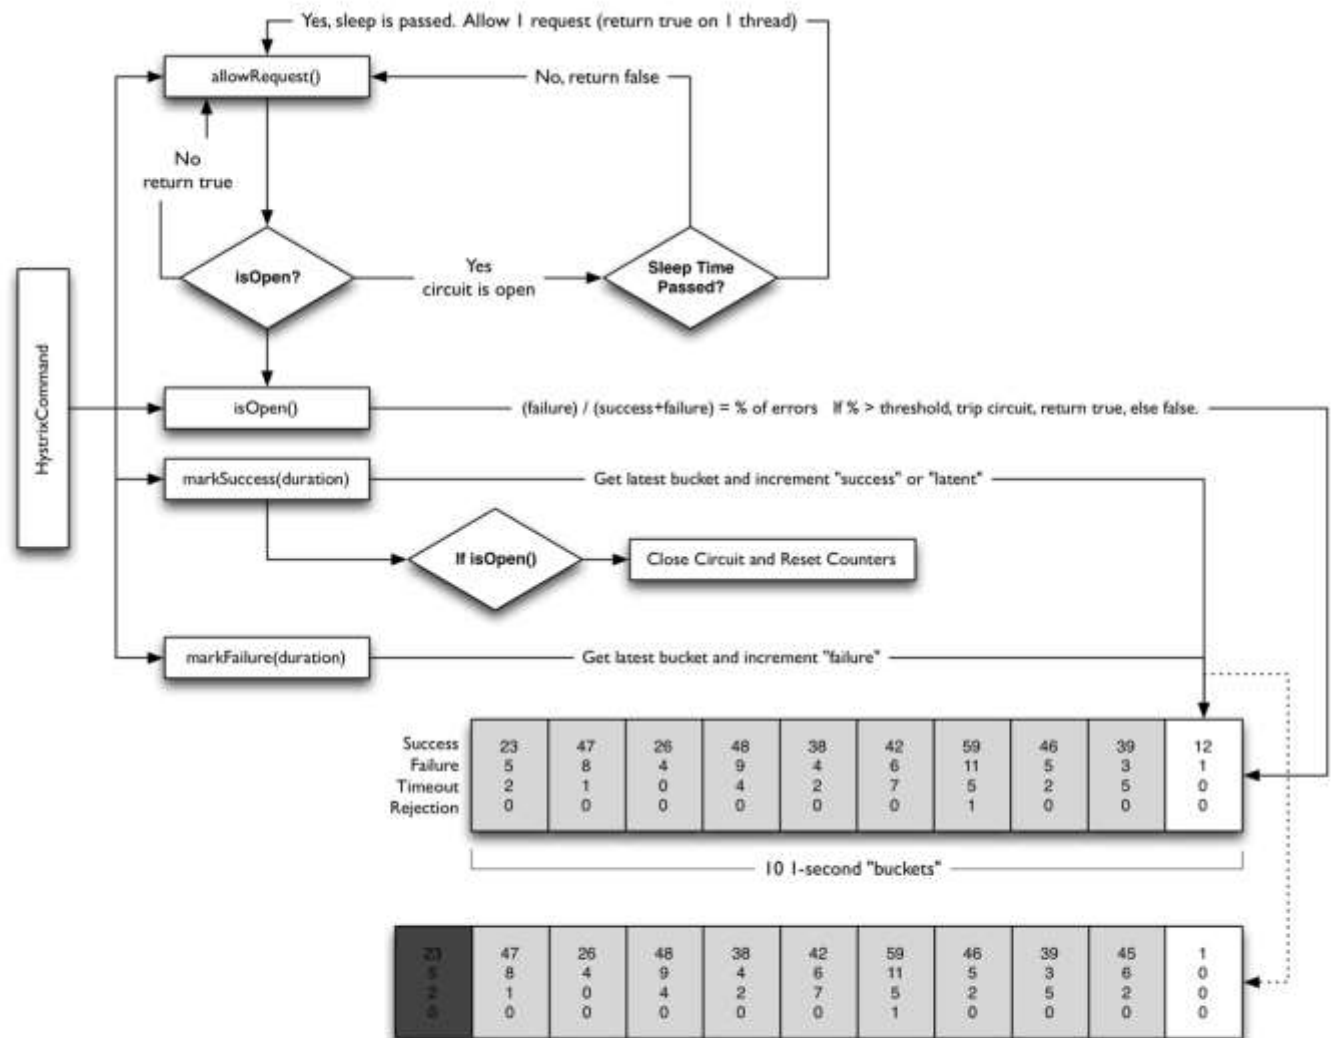
Hystrix Instrumentation

▶ /health

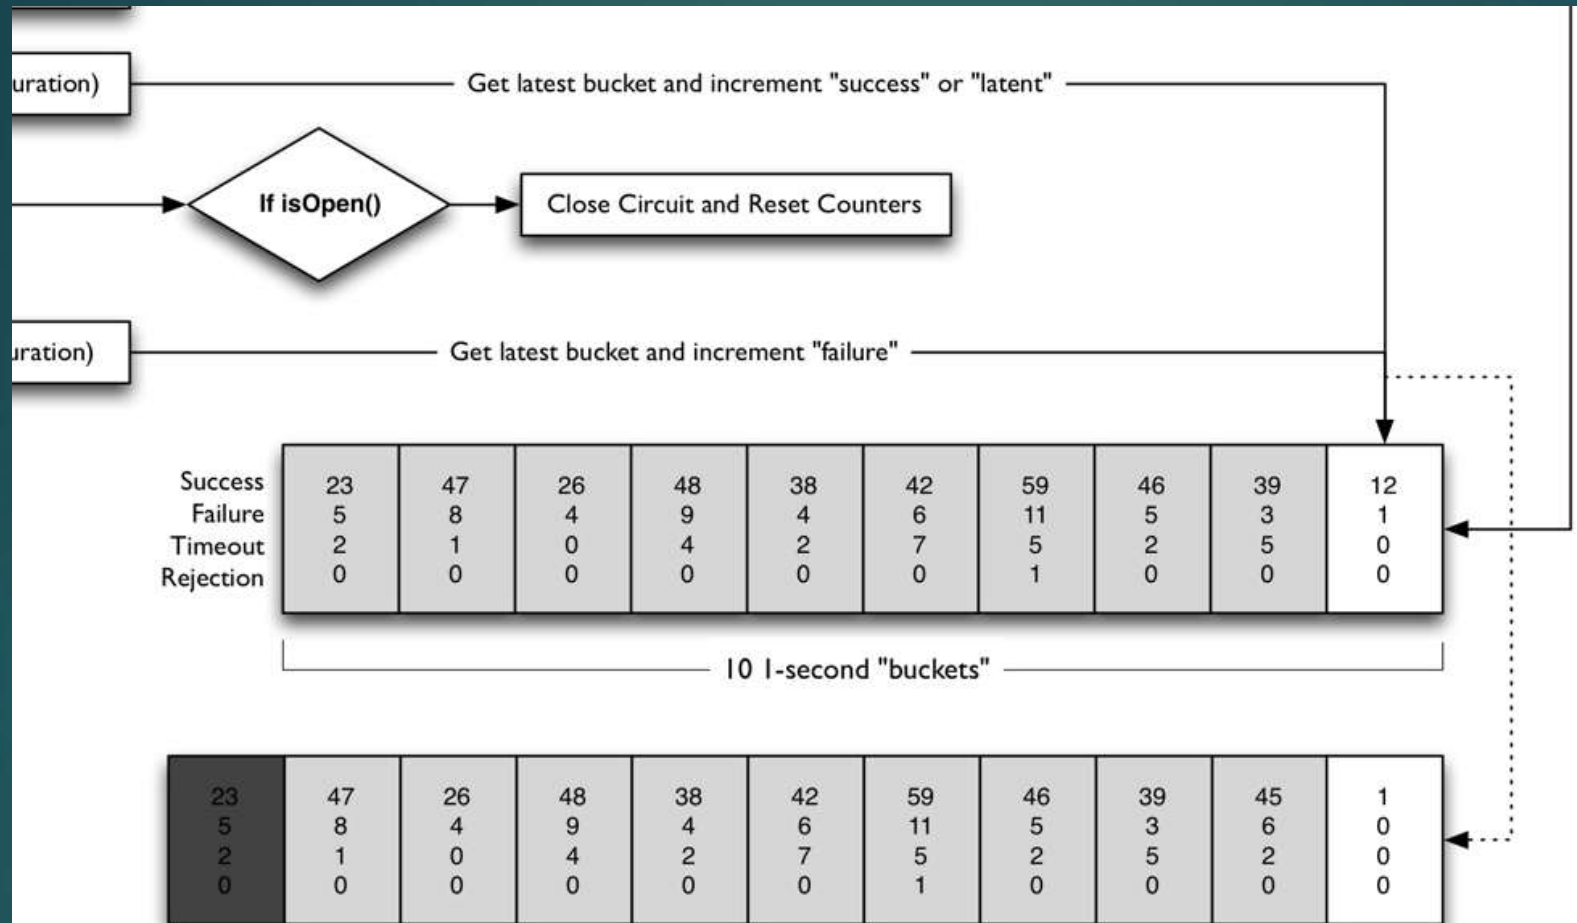
```
{  
  "hystrix": {  
    "openCircuitBreakers": [  
      "StoreServiceGateway::getStoresByZipCode"  
    ],  
    "status": "CIRCUIT_OPEN"  
  },  
  "status": "UP"  
}
```

▶ /hystrix.stream

▶ /metrics



On "getLatestBucket" if the 1-second window is passed a new bucket is created, the rest slid over and the oldest one dropped.



On "getLatestBucket" if the 1-second window is passed a new bucket is created, the rest slid over and the oldest one dropped.

circle color and size represent health and traffic volume

2 minutes of request rate to show relative changes in traffic

hosts reporting from cluster

Hosts
Median
Mean

SubscriberGetAccount

200,545

0

19

94

0 %

0

0

Host: **54.0/s**

Cluster: **20,056.0/s**

Circuit **Closed**

370

1ms

4ms

90th

99th

99.5th

10ms

44ms

61ms

Error percentage of last 10 seconds

Request rate

Circuit-breaker status

last minute latency percentiles

Rolling 10 second counters
with 1 second granularity

Successes	200,545	19	Thread timeouts
Short-circuited (rejected)	0	94	Thread-pool Rejections
		0	Failures/Exceptions

Instrumentation

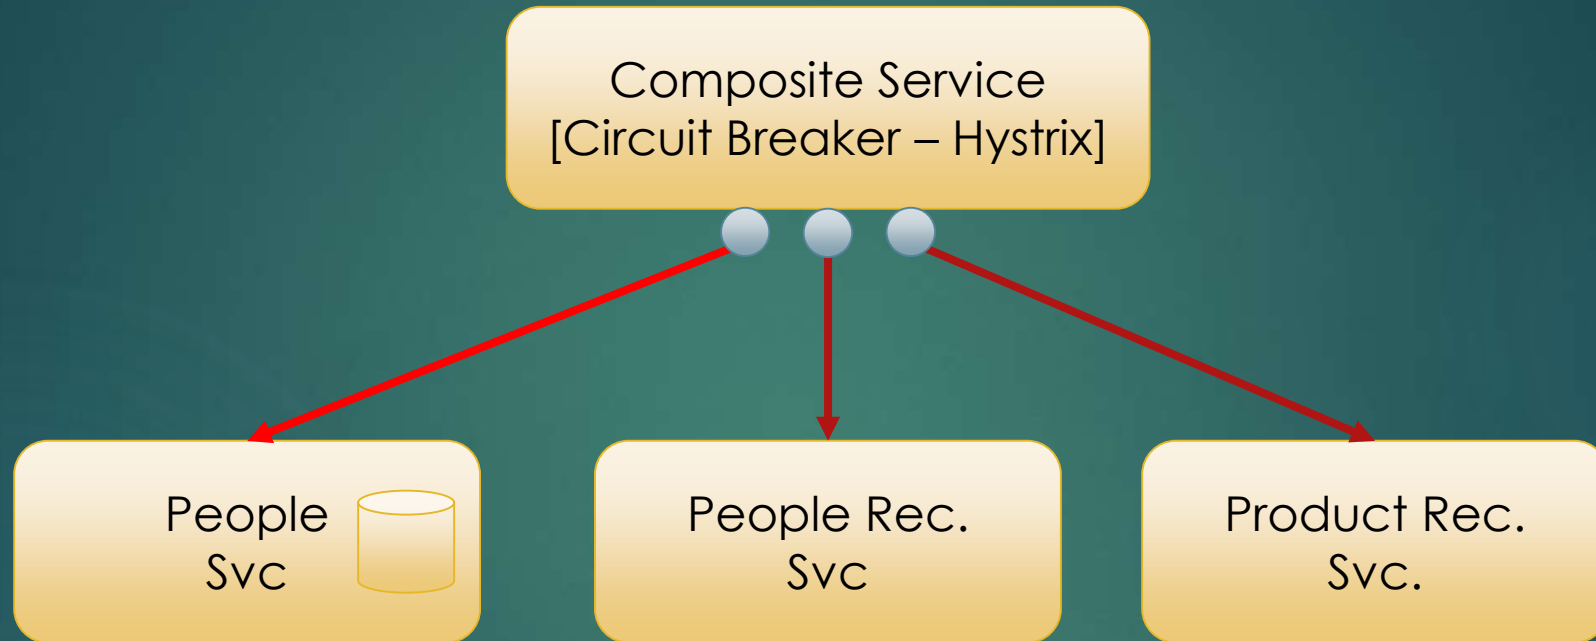
“Measure what is measurable, and make measurable what is not so”.

— Galileo Galilei

“You can’t control what you can’t measure”.

— Tom DeMarco

Controlling Software Projects, Management Measurement & Estimation, (1982)



- Circuit breaker with fallback



- Containerized Spring Boot app with REST API



```
@Service
```

```
public class PersonRecommendationServiceGateway {
```

```
    @Autowired
```

```
    private PersonRecommendationService personRecommendationService;
```

```
    public ResponseEntity<Recommendation[]> getPersonRecommendations(int personId) {  
        return personRecommendationService.getRecommendations(personId);  
    }
```

```
}
```



```
@Service
```

```
public class PersonRecommendationServiceGateway {
```

```
    @Autowired
```

```
    private PersonRecommendationService personRecommendationService;
```

```
    @HystrixCommand
```

```
    public ResponseEntity<Recommendation[]> getPersonRecommendations(int personId) {
```

```
        return personRecommendationService.getRecommendations(personId);
```

```
    }
```

```
}
```



```
@Service
```

```
public class PersonRecommendationServiceGateway {
```

```
    @Autowired
```

```
    private PersonRecommendationService personRecommendationService;
```

```
    @HystrixCommand(fallbackMethod = "defaultPersonRecommendations")
```

```
    public ResponseEntity<Recommendation[]> getPersonRecommendations(int personId) {
```

```
        return personRecommendationService.getRecommendations(personId);
```

```
    }
```

```
    public ResponseEntity<Recommendation[]> defaultPersonRecommendations(int personId) {
```

```
        Recommendation[] emptyArray = {};
```

```
        return new ResponseEntity<Recommendation[]>(emptyArray, HttpStatus.BAD_GATEWAY);
```

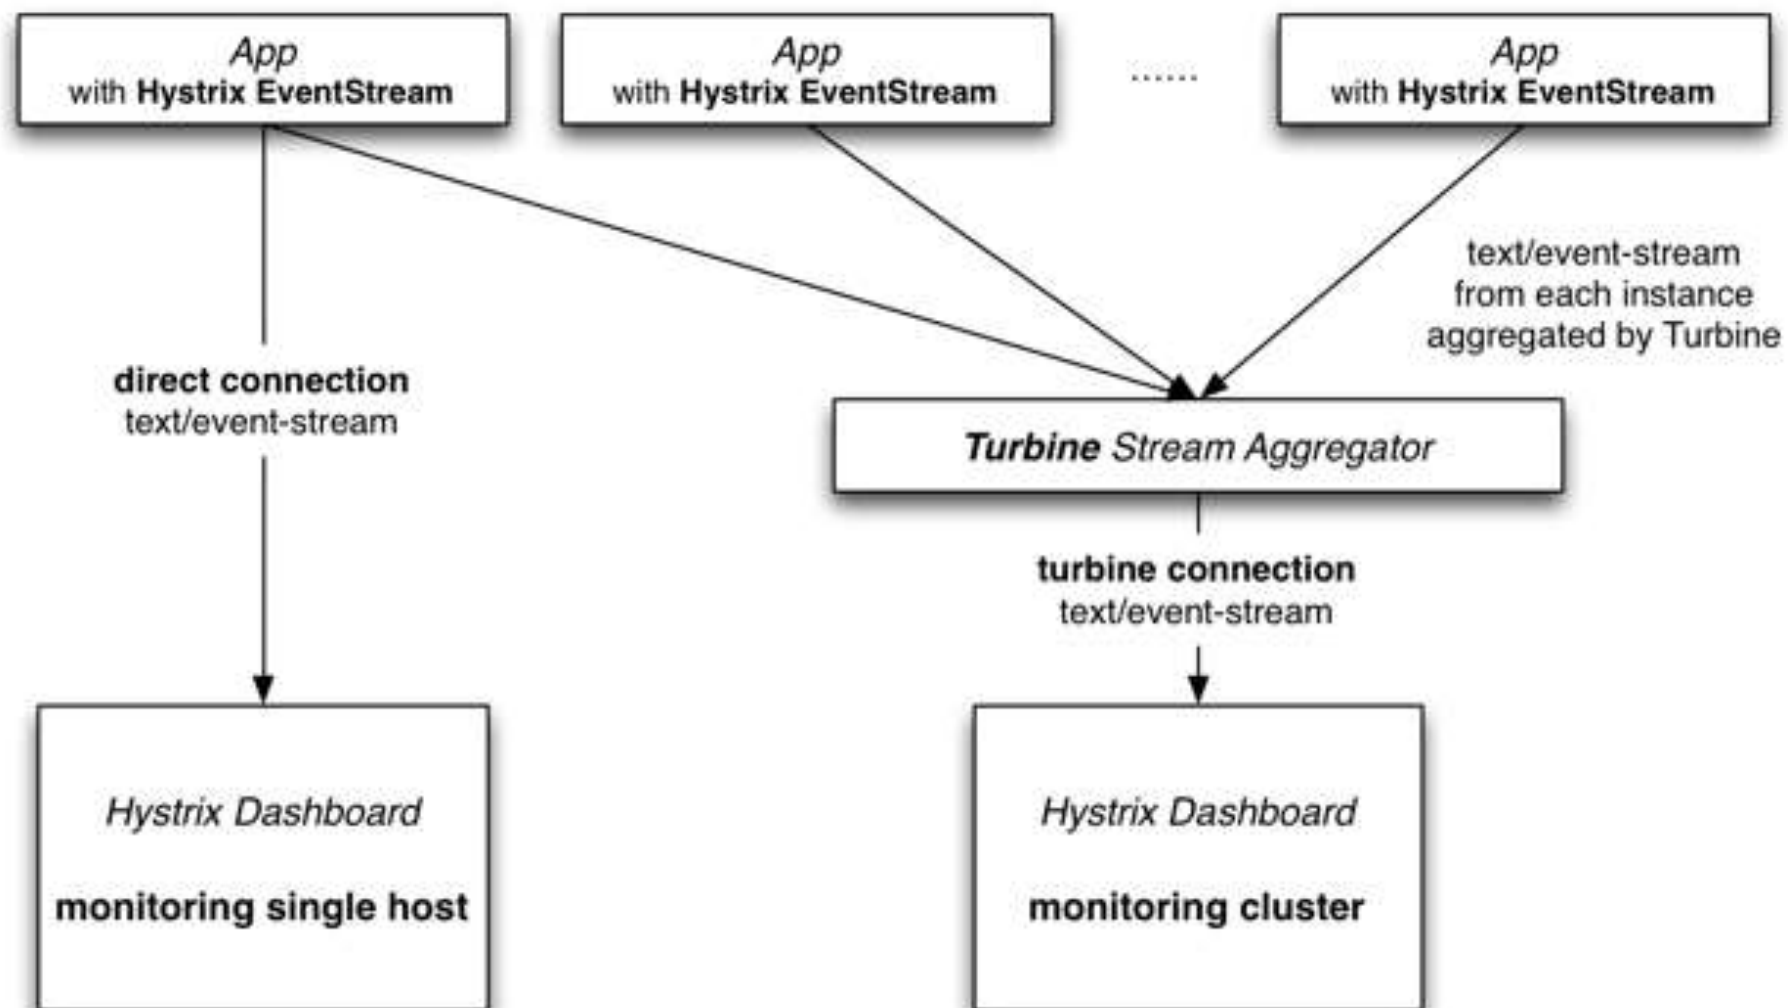
```
    }
```

```
}
```

Demo time!

Recap

- ▶ Why do I care?
 - ▶ Not using cloud!
 - ▶ Not using Spring!
 - ▶ Not using Java!

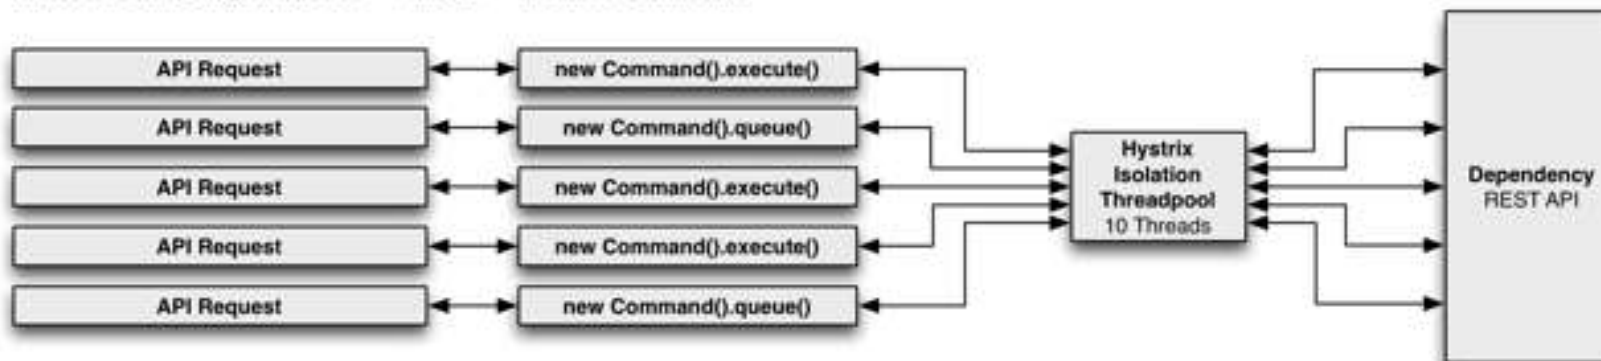


But wait ... there's more

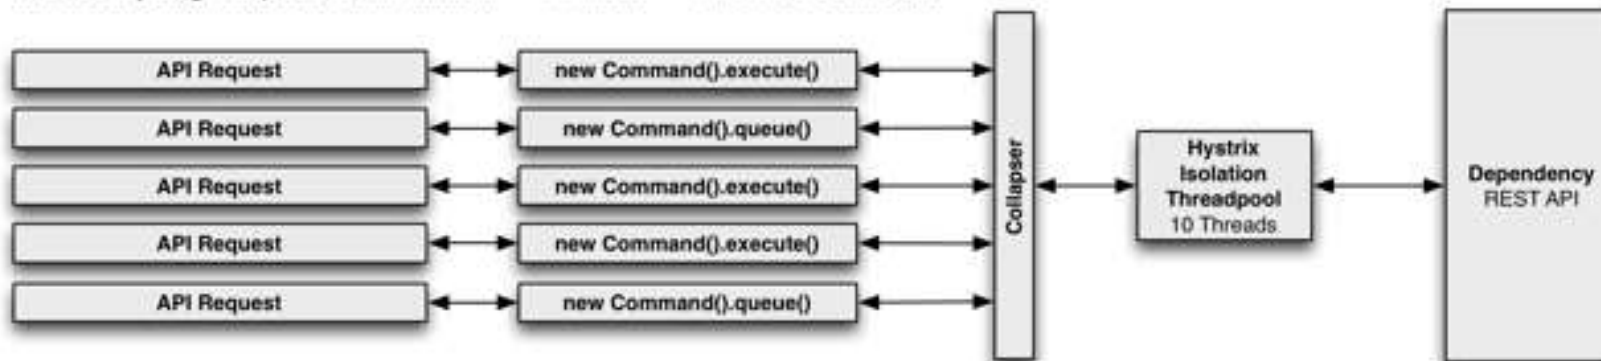
- ▶ Plugin architecture
 - ▶ Event notification
 - ▶ Metrics publication
 - ▶ Property strategy
 - ▶ Concurrency strategy
 - ▶ Command hook execution
- ▶ Additional supported patterns
 - ▶ Request caching
 - ▶ Rate limiting
 - ▶ Request collapsing

Request Collapsing

Without Collapsing: Request == Thread == Network Connection



With Collapsing: Requests within 'window' == 1 Thread == 1 Network Connection



Request Collapsing (cont...)

```
@HystrixCollapser(batchMethod = "getUserByIds")
public Future<User> getUserById(String id) {
    return null;
}
```

```
@HystrixCommand
public List<User> getUserByIds(List<String> ids) {
    List<User> users = new ArrayList<User>();
    for (String id : ids) {
        users.add(new User(id, "name: " + id));
    }
    return users;
}
```

```
Future<User> f1 = userService.getUserById("1");
Future<User> f2 = userService.getUserById("2");
Future<User> f3 = userService.getUserById("3");
Future<User> f4 = userService.getUserById("4");
Future<User> f5 = userService.getUserById("5");
```

Local container efficiencies

- ▶ Docker image size
 - ▶ Java:8 642.4 MB
 - ▶ Java:openjdk-8-jre 310.5 MB
 - ▶ anapsix/alpine-java:jre8 122 MB
(x8 + layers adds up... 4gb not including containers)
- ▶ JVM defaults
 - ▶ -Xmx32m -Xss256k (Memory pool max & thread stack max)
- ▶ Monitoring
 - ▶ Docker stats
 - ▶ cadvisor/ scope

References

- ▶ Hystrix - <https://github.com/Netflix/Hystrix/wiki>
- ▶ Hystrix-Javanica - <https://github.com/Netflix/Hystrix/tree/master/hystrix-contrib/hystrix-javanica>
- ▶ Spring Cloud Netflix - <http://cloud.spring.io/spring-cloud-netflix/spring-cloud-netflix.html>
- ▶ Code – <https://github.com/singram/spring-cloud-microservices>
- ▶ 200 + Circuit breakers on github
<https://github.com/search?utf8=%E2%9C%93&q=circuit+breaker&type=Repositories&ref=searchresults>



Thank You!

@arethoseclams

Stuartingram.com

<https://github.com/singram/spring-cloud-microservices>

Questions?